

# APPLIED MACHINE LEARNING AND EFFICIENT MODEL SELECTION WITH MLR

Bernd Bischl and Michel Lang

UseR 2015, Aalborg

# WELCOME!

- Project home page  
<https://github.com/berndbischl/mlr>
  - ▶ R documentation rendered in HTML
  - ▶ Tutorial for online viewing / download, including many examples
  - ▶ Don't hesitate to interrupt us
  - ▶ There will be a coffee break
- If you do not have `mlr` installed yet, please do so (see wiki page)

# OVERVIEW

INTRODUCTION

WHY MLR?

BUILDING BLOCKS

BENCHMARKING AND MODEL COMPARISON

HYPERPARAMETER TUNING

FEATURE SELECTION

MLR LEARNER WRAPPERS

PARALLELIZATION

VISUALIZATIONS

CARET VS. MLR

OPENML

THE END

# Section 1

## INTRODUCTION


## WHAT IS (SUPERVISED) MACHINE LEARNING?

- Learning structure in data
- The art of predicting stuff
- Model optimization
- Understanding of grey-box models

## DISCLAIMER


- The list is subjective and naively tailored to this talk
- ML is based on math and statistics, we will (mainly) talk about structure, software, and practical issues here

# SUPERVISED CLASSIFICATION TASKS




GOAL: Predict a class (or membership probabilities)

# SUPERVISED REGRESSION TASKS



GOAL: Predict a continuous output


# SUPERVISED SURVIVAL TASKS



GOAL: Predict a survival function  $\hat{S}(t)$ , i.e. the probability to survive to time point  $t$

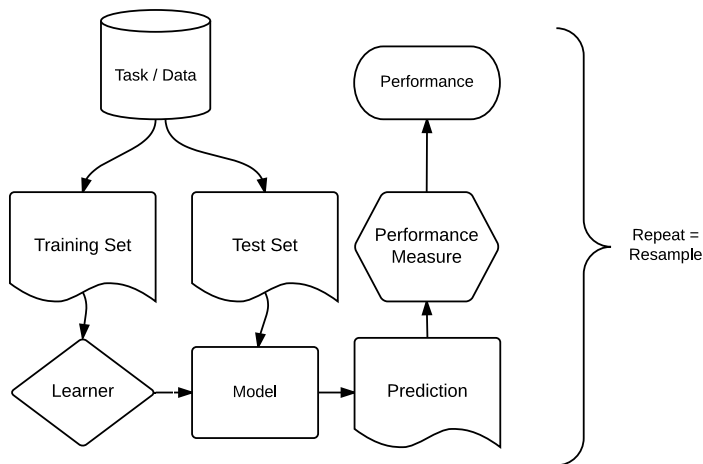


# UNSUPERVISED CLUSTER TASKS



GOAL: Group data into similar clusters (or estimate fuzzy membership probabilities)

# ABSTRACTIONS FOR MACHINE LEARNING



- All learning tasks fit well into this scheme
- Challenge as statistician: find a suitable model to maximize the outcome (or minimize the loss)

## Section 2

WHY MLR?

# MOTIVATION

## THE GOOD NEWS

- CRAN serves hundreds of packages for machine learning (cf. CRAN task view machine learning)
- Many packages are compliant to the unwritten interface definition:

```
> model = fit(target ~ ., data = train.data, ...)  
> predictions = predict(model, newdata = test.data, ...)
```

# MOTIVATION

## THE BAD NEWS

- Some packages do not support the formula interface or their API is “just different”
- Functionality is always package or model-dependent, even though the procedure might be general
- No meta-information available or buried in docs (sometimes not documented at all)
- Many packages require the user to “guess” good hyperparameters
- Larger experiments lead to lengthy, tedious and error-prone code

Our goal: A domain-specific language for many machine learning concepts!


# MOTIVATION: MLR

- Unified interface for the basic building blocks: tasks, learners, resampling, hyperparameters, ...
- Reflections: nearly all objects are queryable (i.e. you can ask them for their properties and program on them)
- The OO-structure allows many generic algorithms:
  - ▶ Bagging
  - ▶ Stacking
  - ▶ Feature Selection
  - ▶ ...
- Easily extensible via S3
  - ▶ Extension is not covered here, but explained in detail in the online tutorial
  - ▶ You do not need to understand S3 to use `mlr`
  - ▶ Wondering why we don't use S4? We care about code bloat and speed.

## Section 3

# BUILDING BLOCKS

# BUILDING BLOCKS



- `mlr` objects: tasks, learners, measures, resampling instances.



# TASK ABSTRACTION

- Tasks encapsulate data and meta-information about it
- Regression, classification, clustering, survival tasks
- Data is stored inside an environment to save memory

```
> task = makeClassifTask(data = iris, target = "Species")  
> print(task)
```

```
## Supervised task: iris  
## Type: classif  
## Target: Species  
## Observations: 150  
## Features:  
## numerics  factors  ordered  
##          4         0         0  
## Missings: FALSE  
## Has weights: FALSE  
## Has blocking: FALSE  
## Classes: 3  
##      setosa versicolor  virginica  
##      50         50         50  
## Positive class: NA
```

# TASK ABSTRACTION: API I

```
> getTaskId(task)

## [1] "iris"

> str(getTaskData(task))

## 'data.frame': 150 obs. of  5 variables:
##  $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
##  $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
##  $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
##  $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
##  $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1
```

# TASK ABSTRACTION: API II

```
> str(getTaskDescription(task))

## List of 11
## $ id          : chr "iris"
## $ type        : chr "classif"
## $ target      : chr "Species"
## $ size        : int 150
## $ n.feats     : Named int [1:3] 4 0 0
## $ ..- attr(*, "names")= chr [1:3] "numerics" "factors" "ordered"
## $ has.missings: logi FALSE
## $ has.weights : logi FALSE
## $ has.blocking: logi FALSE
## $ class.levels: chr [1:3] "setosa" "versicolor" "virginica"
## $ positive    : chr NA
## $ negative    : chr NA
## $ - attr(*, "class")= chr [1:2] "TaskDescClassif" "TaskDesc"
```

# TASK ABSTRACTION: API III

```
> getTaskSize(task)

## [1] 150

> getTaskFeatureNames(task)

## [1] "Sepal.Length" "Sepal.Width" "Petal.Length" "Petal.Width"

> getTaskTargetNames(task)

## [1] "Species"

> getTaskFormula(task)

## Species ~ .
## <environment: 0xd911ed8>

> summary(getTaskTargets(task))

##      setosa versicolor  virginica
##      50         50         50
```

# LEARNER ABSTRACTION I

- Internal structure of learners:
  - ▶ wrappers around `fit()` and `predict()` of the package
  - ▶ description of the parameter set
  - ▶ annotations
- Naming convention: `<tasktype>.<functionname>`

```
> makeLearner("classif.rpart")  
> makeLearner("regr.rpart")
```

- Adding custom learners is covered in the tutorial

# LEARNER ABSTRACTION II

```
> lrn = makeLearner("classif.rpart")
> print(lrn)

## Learner classif.rpart from package rpart
## Type: classif
## Name: Decision Tree; Short name: rpart
## Class: classif.rpart
## Properties: twoclass,multiclass,missings,numerics,factors,ordered,prob,weigh
## Predict-Type: response
## Hyperparameters: xval=0
```

# WHAT LEARNERS ARE AVAILABLE? I

## CLASSIFICATION (54)

- LDA, QDA, RDA, MDA
- Trees and forests
- Boosting (different variants)
- SVMs (different variants)
- ...

## CLUSTERING (6)

- K-Means
- EM
- DBscan
- X-Means
- ...

## REGRESSION (45)

- Linear, lasso and ridge
- Boosting
- Trees and forests
- Gaussian processes
- ...

## SURVIVAL (10)

- Cox-PH
- Cox-Boost
- Random survival forest
- Penalized regression
- ...

# WHAT LEARNERS ARE AVAILABLE? II

We can explore them on the webpage – or ask `mlr`

```
> # list all classification learners which can predict probabilities
> # and allow multiclass classification
> head(unname(
+   listLearners("classif", properties = c("prob", "multiclass"))
+ ))

## [1] "classif.bdk"          "classif.boosting"    "classif.cforest"
## [4] "classif.ctree"       "classif.extraTrees"  "classif.gbm"
```



# WHAT LEARNERS ARE AVAILABLE? III

Get all applicable learners for a task

```
> head(unname(listLearners(task)))  
  
## [1] "classif.bdk"          "classif.boosting"    "classif.cforest"  
## [4] "classif.ctree"        "classif.extraTrees"  "classif.fnn"
```

# PARAMETER ABSTRACTION

- Extensive meta-information for hyperparameters available: storage type, constraints, defaults, dependencies
- Automatically checked for feasibility
- You can program on parameters!

```
> getParamSet(lrn)
```

##	Type	len	Def	Constr	Req	Trafo
## minsplit	integer	-	20	1 to Inf	-	-
## minbucket	integer	-	-	1 to Inf	-	-
## cp	numeric	-	0.01	0 to 1	-	-
## maxcompete	integer	-	4	0 to Inf	-	-
## maxsurrogate	integer	-	5	0 to Inf	-	-
## usesurrogate	discrete	-	2	0,1,2	-	-
## surrogatestyle	discrete	-	0	0,1	-	-
## maxdepth	integer	-	30	1 to 30	-	-
## xval	integer	-	10	0 to Inf	-	-
## parms	untyped	-	-	-	-	-

# LEARNER ABSTRACTION: API

```
> lrn$properties

## [1] "twoclass"      "multiclass" "missings"    "numerics"    "factors"
## [6] "ordered"      "prob"       "weights"

> getHyperPars(lrn)

## $xval
## [1] 0

> lrn = setHyperPars(lrn, cp = 0.3)
> lrn = setPredictType(lrn, "prob")
> lrn = setPredictThreshold(lrn, 0.7);
```

# PERFORMANCE MEASURES

- Performance measures evaluate the predictions a test set and aggregate them over multiple in resampling iterations
- 22 classification, 7 regression, 5 cluster, 1 survival
- Internally: performance function, default aggregation function and annotations
- Adding custom measures is covered in the tutorial

```
> print(mmce)

## Name: Mean misclassification error
## Performance measure: mmce
## Properties: classif,classif.multi,req.pred,req.truth
## Minimize: TRUE
## Best: 0; Worst: 1
## Aggregated by: test.mean
## Note:
```

# WHAT MEASURES ARE AVAILABLE?

We can explore them on the webpage – or ask mlr

```
> listMeasures("classif")
```

```
## [1] "f1"           "featperc"      "mmce"
## [4] "tn"           "tp"            "mcc"
## [7] "fn"           "fp"            "npv"
## [10] "bac"          "timeboth"      "acc"
## [13] "ppv"          "multiclass.auc" "brier"
## [16] "fnr"          "auc"           "tnr"
## [19] "ber"          "timepredict"   "fpr"
## [22] "gmean"        "tpr"           "gpr"
## [25] "fdr"          "timetrain"
```

```
> listMeasures(task)
```

```
## [1] "featperc"      "mmce"           "timeboth"
## [4] "acc"           "multiclass.auc" "ber"
## [7] "timepredict"   "timetrain"
```

## R EXAMPLE

Training and prediction

# RESAMPLING ABSTRACTION I

- Procedure: Train, Predict, Eval, Repeat.
- Aim: Estimate expected model performance.
  - ▶ Hold-Out
  - ▶ Cross-validation (normal, repeated)
  - ▶ Bootstrap (OOB, B632, B632+)
  - ▶ Subsampling
  - ▶ Stratification
  - ▶ Blocking
- Instantiate it or not (= create data split indices)

```
> rdesc = makeResampleDesc("CV", iters = 3)
> rin = makeResampleInstance(rdesc, task = task)
> str(rin$train.inds)
```

```
## List of 3
## $ : int [1:100] 68 75 78 11 31 145 63 110 123 9 ...
## $ : int [1:100] 68 75 27 11 41 128 63 110 58 98 ...
## $ : int [1:100] 27 78 41 31 145 128 123 58 9 46 ...
```

# RESAMPLING ABSTRACTION II

## RESAMPLING A LEARNER

- Measures on test (or train) sets
- Returns aggregated values, predictions and some useful extra information

```
> lrn = makeLearner("classif.rpart")  
> rdesc = makeResampleDesc("CV", iters = 3)  
> measures = list(mmce, timetrain)  
> r = resample(lrn, task, rdesc, measures = measures)
```

- For the lazy

```
> r = crossval(lrn, task, iters = 3, measures = measures)
```



# RESAMPLING ABSTRACTION III

```
> print(r)

## Resample Result
## Task: iris
## Learner: classif.rpart
## mmce.aggr: 0.05
## mmce.mean: 0.05
## mmce.sd: 0.03
## timetrain.aggr: 0.00
## timetrain.mean: 0.00
## timetrain.sd: 0.00
## Runtime: 0.0250013
```

# RESAMPLING ABSTRACTION IV

```
> names(r)

## [1] "learner.id"      "task.id"          "measures.train"
## [4] "measures.test"   "aggr"             "pred"
## [7] "models"          "err.msgs"         "extract"
## [10] "runtime"

> r$measures.test

##   iter    mmce timetrain
## 1     1 0.02667     0.005
## 2     2 0.06667     0.004

> r$aggr

##           mmce.test.mean timetrain.test.mean
##                0.04667                0.00450
```

# RESAMPLING ABSTRACTION V

```
> head(as.data.frame(r$pred))
```

##	id	truth	response	iter	set
## 3	3	setosa	setosa	1	test
## 5	5	setosa	setosa	1	test
## 8	8	setosa	setosa	1	test
## 11	11	setosa	setosa	1	test
## 13	13	setosa	setosa	1	test
## 15	15	setosa	setosa	1	test

# CONFIGURING THE PACKAGE

- What to do when training fails? error, warn, or be quiet?
  - You don't want to stop in complex loops like benchmark
  - `FailureModel` is created that predicts NAs
- Show verbose info messages?
- What if parameters are not described in learner?
- `?configureMlr` sets global flags and can be overwritten for individual learners

## Section 4

# BENCHMARKING AND MODEL COMPARISON

# BENCHMARKING AND MODEL COMPARISON I

## BENCHMARKING

- Comparison of multiple models on multiple data sets
- Aim: Find best learners for a data set or domain, learn about learner characteristics, ...

# BENCHMARKING AND MODEL COMPARISON II

## BENCHMARKING IN MLR

- Train and test sets are synchronized, i.e. all learners see the same data splits
- Can be done in parallel (see later)
- Can be combined with feature selection / tuning / nested resampling (see later)
- Results stored in well-defined container object, with getters and converters
- We are working on standard analysis tools

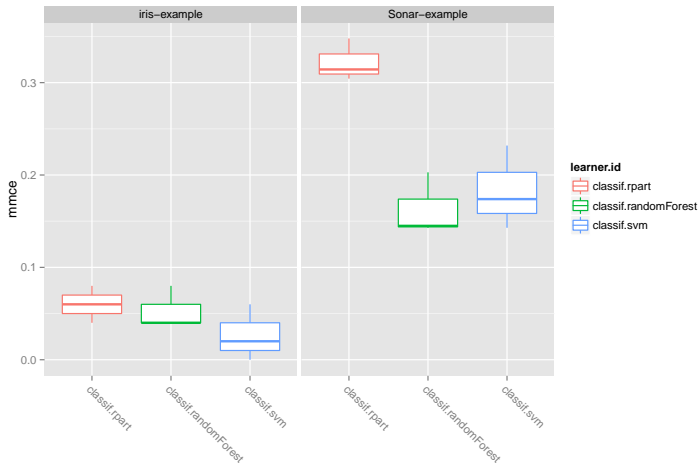
# BENCHMARKING AND MODEL COMPARISON III

```
> library(mlr)
> # lets try a couple of methods on some (mlr example) tasks
>
> # these are predefined in mlr for toying around:
> tasks = list(iris.task, sonar.task)
>
> learners = list(
+   makeLearner("classif.rpart"),
+   makeLearner("classif.randomForest", ntree = 500),
+   makeLearner("classif.svm")
+ )
>
> rdesc = makeResampleDesc("CV", iters = 3)
> set.seed(1)
> br = benchmark(learners, tasks, rdesc)
```



# BENCHMARKING AND MODEL COMPARISON IV

```
> plotBenchmarkResult(br)
```



# BENCHMARKING AND MODEL COMPARISON V

```
> getBMRAggrPerformances(br, as.df = TRUE)
```

##	task.id	learner.id	mmce.test.mean
## 1	iris-example	classif.rpart	0.06000
## 2	iris-example	classif.randomForest	0.05333
## 3	iris-example	classif.svm	0.02667
## 4	Sonar-example	classif.rpart	0.32215
## 5	Sonar-example	classif.randomForest	0.16356
## 6	Sonar-example	classif.svm	0.18288

# BENCHMARKING AND MODEL COMPARISON VI

```
> getBMRPerformances(br, as.df = TRUE)
```

##	task.id	learner.id	iter	mmce
## 1	iris-example	classif.rpart	1	0.0600
## 2	iris-example	classif.rpart	2	0.0400
## 3	iris-example	classif.rpart	3	0.0800
## 4	iris-example	classif.randomForest	1	0.0400
## 5	iris-example	classif.randomForest	2	0.0400
## 6	iris-example	classif.randomForest	3	0.0800
## 7	iris-example	classif.svm	1	0.0000
## 8	iris-example	classif.svm	2	0.0200
## 9	iris-example	classif.svm	3	0.0600
## 10	Sonar-example	classif.rpart	1	0.3478
## 11	Sonar-example	classif.rpart	2	0.3043
## 12	Sonar-example	classif.rpart	3	0.3143
## 13	Sonar-example	classif.randomForest	1	0.2029
## 14	Sonar-example	classif.randomForest	2	0.1449
## 15	Sonar-example	classif.randomForest	3	0.1429
## 16	Sonar-example	classif.svm	1	0.2319
## 17	Sonar-example	classif.svm	2	0.1739
## 18	Sonar-example	classif.svm	3	0.1429

# BENCHMARKING AND MODEL COMPARISON VII

```
> head(getBMRPredictions(br, as.df = TRUE), 10)
```

##	task.id	learner.id	id	truth	response	iter	set
## 1	iris-example	classif.rpart	1	setosa	setosa	1	test
## 2	iris-example	classif.rpart	3	setosa	setosa	1	test
## 3	iris-example	classif.rpart	12	setosa	setosa	1	test
## 4	iris-example	classif.rpart	17	setosa	setosa	1	test
## 5	iris-example	classif.rpart	22	setosa	setosa	1	test
## 6	iris-example	classif.rpart	24	setosa	setosa	1	test
## 7	iris-example	classif.rpart	25	setosa	setosa	1	test
## 8	iris-example	classif.rpart	26	setosa	setosa	1	test
## 9	iris-example	classif.rpart	31	setosa	setosa	1	test
## 10	iris-example	classif.rpart	34	setosa	setosa	1	test

## Section 5

# HYPERPARAMETER TUNING

# HYPERPARAMETER TUNING I

## TUNING

- Used to find “best” hyperparameters for a method in a data-dependent way
- Essential for some methods, e.g. SVMs

## TUNING IN MLR

- General procedure: Tuner proposes param point, eval by resampling, feedback value to tuner
- Multiple tuners through exactly the same interface
- All evals and more info is logged into `OptPath` object

# HYPERPARAMETER TUNING II

## GRID SEARCH

- Basic method: Exhaustively try all combinations of finite grid
- Inefficient, combinatorial explosion
- Searches large, irrelevant areas
- Reasonable for continuous parameters?
- Still often default method

## RANDOM SEARCH

- Randomly draw parameters
- `mlr` supports all types and dependencies
- Scales better than grid search, easily extensible

# R EXAMPLE

Tuning




# AUTOMATIC MODEL SELECTION

## PRIOR APPROACHES:


- Looking for the silver bullet model
  - ~ Failure
- Exhaustive benchmarking / search
  - ~ Per data set: too expensive
  - ~ Over many: contradicting results
- Meta-Learning:
  - ~ Failure
  - ~ Usually not for preprocessing / hyperparameters

GOAL: Data dependent + Automatic + Efficient


# BLACK-BOX-PERSPECTIVE IN CONFIGURATION




# BLACK-BOX-PERSPECTIVE IN CONFIGURATION




# BLACK-BOX-PERSPECTIVE IN CONFIGURATION



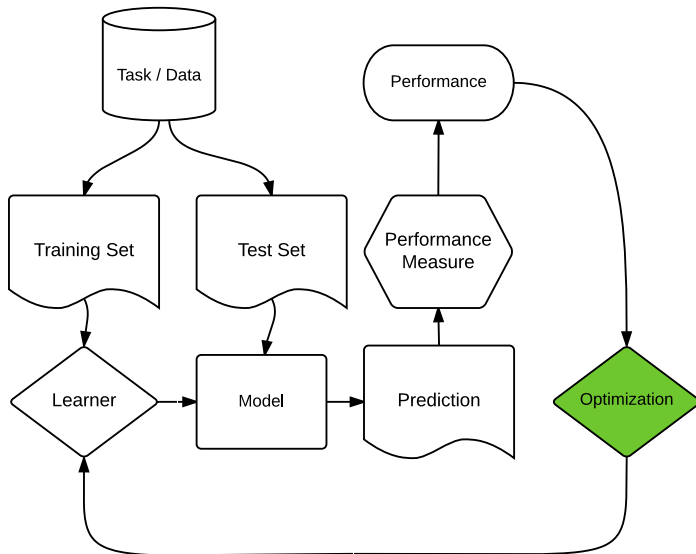
# BLACK-BOX-PERSPECTIVE IN CONFIGURATION



# BLACK-BOX-PERSPECTIVE IN CONFIGURATION



# ADAPTIVE TUNING




# GENERAL ALGORITHM CONFIGURATION

- Assume a (parametrized) algorithm  $a$
- Parameter space  $\theta \in \Theta$   
might be discrete and dependent / hierarchical
- Stochastic generating process for instances  $i \sim P$ , where we draw i.i.d. from.
- Run algorithm  $a$  on  $i$  and measure performance  
 $f(i, \theta) = \text{run}(i, a(\theta))$
- Objective:  $\min_{\theta \in \Theta} E_P[f(i, \theta)]$
- No derivative for  $f(\cdot, \theta)$ , black-box
- $f$  is stochastic / noisy
- $f$  is likely expensive to evaluate
- Consequence: very hard problem

$\leadsto$  RACING OR MODEL-BASED / BAYESIAN OPTIMIZATION




# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING



- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.


# IDEA OF (F-)RACING



- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.




# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING




- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING



- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.


# IDEA OF (F-)RACING



- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.




# IDEA OF (F-)RACING



- Write down all candidate solutions
- Iterate the following till budget exhausted
  - One “generation”
    - ▶ Evaluate all candidates on an instance, and another, . . .
    - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
    - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF (F-)RACING



- Write down all candidate solutions
- Iterate the following till budget exhausted
- One “generation”
  - ▶ Evaluate all candidates on an instance, and another, . . .
  - ▶ After some time, compare candidates via statistical test, e.g., Friedman test with post-hoc analysis for pairs
  - ▶ Remove outperformed candidates
- Output: Remaining candidates
- Yes, the testing completely ignores “sequentiality” and is somewhat heuristic.

# IDEA OF ITERATED F-RACING

## WHAT MIGHT BE PROBLEMATIC?

- We might have many or an infinite number of candidates

## ITERATED RACING

- Have a stochastic model to draw candidates from in every generation
- For each parameter: Univariate, independent distribution (factorized joint distribution)
- Sample distributions centered at “elite” candidates from previous generation(s)
- Reduce distributions’ width / variance in later generations for convergence

# IDEA OF ITERATED F-RACING

## WHATS GOOD ABOUT THIS

- Very simple and generic algorithm
- Can easily be parallelized
- A nice R package exists: `irace`<sup>1</sup>

## WHAT MIGHT BE NOT SO GOOD

- Quite strong (wrong?) assumptions in the probability model
- Sequential model-based optimization is probably more efficient  
(But be careful: Somewhat my personal experience and bias,  
as not so many large scale comparisons exist)

---

<sup>1</sup>Lopez-Ibanez et al, “The irace package, Iterated Race for Automatic Algorithm Configuration. Technical Report TR/IRIDIA/2011-004, IRIDIA, Université libre de Bruxelles, Belgium, 2011.”

## Section 6

# FEATURE SELECTION

# FEATURE SELECTION I

- Reduce dimensionality, increase interpretability and predictive performance
- Concepts:
  - FILTER: Preliminary step, independent from model
  - WRAPPER: Wrapped around model fit which is iteratively scored
  - EMBEDDED: Model has feature selection embedded, e.g. lasso regression

# FEATURE SELECTION II

## FEATURE FILTERS

- Usually: Quickly compute a numerical score per feature
- Encodes influence of feature on output
- Often independent of ML model
- Often fast to compute
- Can be used to visualize data structure
- Can be used to rank or threshold the feature set, and to reduce feature set size
- Terrible if complex correlations exist

# FEATURE SELECTION III

## FILTER EXAMPLES

- Correlation between  $x_i$  and  $y$  in regression
- Mutual information in classification
- The random forest importance value
- $\chi^2$ -statistic for independence between  $x_i$  and  $y$



# FEATURE SELECTION IV

```
> fv = generateFilterValuesData(iris.task, method = "information.gain")
> print(fv)

## FilterValues:
## Task: iris-example
##           name      type information.gain
## 1 Sepal.Length numeric          0.4521
## 2 Sepal.Width  numeric          0.2673
## 3 Petal.Length numeric          0.9403
## 4 Petal.Width  numeric          0.9554


> task2 = filterFeatures(iris.task, fval = fv, perc = 0.5)
> print(getTaskFeatureNames(task2))

## [1] "Petal.Length" "Petal.Width"
```

You can optimize this selection threshold jointly with the model!

# FEATURE SELECTION V

```
> plotFilterValues(fv)
```




# FEATURE SELECTION VI

## WRAPPER APPROACH

- Evaluate feature sets with learner, e.g. by cross-validation
- Measures probably what you are interested in
- Will be slow in very high-dimensional spaces
- Sequential Forward Search (SFS) (or backward)
- Sequential Floating Forward Search (SFFS)
- Genetic Algorithm (GA)

# FEATURE SELECTION VII



# FEATURE SELECTION VIII

```
> ## Specify the search strategy
> ctrl = makeFeatSelControlSequential(method = "sfs", alpha = 0.05)
>
> ## Select features
> rdsc = makeResampleDesc("CV", iters = 10)
> sfeats = selectFeatures(learner = "regr.lm", task = bh.task,
+   resampling = rdsc, control = ctrl, show.info = FALSE)
> sfeats

## FeatSel result:
## Features (11): crim, zn, chas, nox, rm, dis, rad, tax, ptratio, b, lstat
## mse.test.mean=23.7
```

# FEATURE SELECTION IX

```
> analyzeFeatSelResult(sfeats)
```

```
## Features          : 11
## Performance       : mse.test.mean=23.7
## crim, zn, chas, nox, rm, dis, rad, tax, ptratio, b, lstat
##
## Path to optimum:
## - Features:      0  Init   :                Perf = 84.896  Diff: NA   *
## - Features:      1  Add    : lstat          Perf = 39.013  Diff: 45.883  *
## - Features:      2  Add    : rm              Perf = 31.553  Diff: 7.4592  *
## - Features:      3  Add    : ptratio         Perf = 27.992  Diff: 3.5617  *
## - Features:      4  Add    : dis             Perf = 27.189  Diff: 0.80245 *
## - Features:      5  Add    : nox             Perf = 25.734  Diff: 1.4555  *
## - Features:      6  Add    : b               Perf = 25.207  Diff: 0.52638 *
## - Features:      7  Add    : zn              Perf = 24.935  Diff: 0.27213 *
## - Features:      8  Add    : chas            Perf = 24.73   Diff: 0.20546 *
## - Features:      9  Add    : rad             Perf = 24.595  Diff: 0.1346  *
## - Features:     10  Add    : tax             Perf = 24.11   Diff: 0.48483 *
## - Features:     11  Add    : crim            Perf = 23.695  Diff: 0.41533 *
##
## Stopped, because no improving feature was found.
```

## Section 7

### MLR LEARNER WRAPPERS

# MLR LEARNER WRAPPERS I

## WHAT?

- Extend the functionality of learners by adding an `mlr` wrapper to them
- The wrapper hooks into the `train` and `predict` of the base learner and extends it
- This way, you can create a new `mlr` learner with extended functionality
- Hyperparameter definition spaces get joined!




# MLR LEARNER WRAPPERS II

## AVAILABLE WRAPPERS

- PREPROCESSING: PCA, normalization (z-transformation)
- PARAMETER TUNING: grid, optim, random search, genetic algorithms, CMAES, iRace, MBO
- FILTER: correlation- and entropy-based,  $\chi^2$ -test, mRMR, ...
- FEATURE SELECTION: (floating) sequential forward/backward, exhaustive search, genetic algorithms, ...
- IMPUTE: dummy variables, imputations with mean, median, min, max, empirical distribution or other learners
- BAGGING to fuse learners on bootstrapped samples
- STACKING to combine models in heterogenous ensembles
- OVER- AND UNDERSAMPLING for unbalanced classification

# NESTED RESAMPLING

- Using the TuningWrapper or FeatureSelectionWrapper allows to enable nested resampling
- Ensures **unbiased** results for model optimization
- Everything else is statistically unsound



# R EXAMPLE

Complex tuning example

## Section 8

# PARALLELIZATION

# PARALLELIZATION I

- We use our own package: `parallelMap`
- Initialize a backend with `parallelStart`
- Stop with `parallelStop`

```
> parallelStart("multicore")  
> benchmark(...)  
> parallelStop()
```

- Backends: `local`, `multicore`, `socket`, `mpi` and `BatchJobs`
- The latter means support for: makeshift SSH-clusters and HPC schedulers like SLURM, Torque/PBS, SGE or LSF
- The first loop which is marked as parallel executable will be automatically parallelized

# PARALLELIZATION II

## PARALLELIZATION LEVELS

- Which loop to parallelize depends on number of iterations
- Levels allow fine grained control over the parallelization
  - ▶ `mlr.resample`: Each resampling iteration (a train / test step) is a parallel job.
  - ▶ `mlr.benchmark`: Each experiment “run this learner on this data set” is a parallel job.
  - ▶ `mlr.tuneParams`: Each evaluation in hyperparameter space “resample with these parameter settings” is a parallel job. How many of these can be run independently in parallel depends on the tuning algorithm.
  - ▶ `mlr.selectFeatures`: Each evaluation in feature space “resample with this feature subset” is a parallel job.

# PARALLELIZATION III

```
> lrns = list(makeLearner("classif.rpart"), makeLearner("classif.svm"))
> rdesc = makeResampleDesc("Bootstrap", iters = 100)
>
> parallelStart("multicore", 8)
```

```
## Starting parallelization in mode=multicore with cpus=8.
```

```
> benchmark(learners = lrns, tasks = iris.task, resamplings = rdesc)
```

```
## Mapping in parallel: mode = multicore; cpus = 8; elements = 2.
```

```
##      task.id    learner.id mmce.test.mean
## 1 iris-example classif.rpart      0.05659
## 2 iris-example  classif.svm      0.04218
```

```
> parallelStop()
```

```
## Stopped parallelization. All cleaned up.
```

# PARALLELIZATION IV

Parallelize the bootstrap instead:

```
> parallelStart("multicore", 8, level = "mlr.resample")

## Starting parallelization in mode=multicore with cpus=8.

> benchmark(learners = lrns, tasks = iris.task, resamplings = rdesc)

## Mapping in parallel: mode = multicore; cpus = 8; elements = 100.
## Mapping in parallel: mode = multicore; cpus = 8; elements = 100.

##      task.id    learner.id mmce.test.mean
## 1 iris-example classif.rpart      0.05810
## 2 iris-example  classif.svm       0.04363

> parallelStop()

## Stopped parallelization. All cleaned up.
```



## Section 9

# VISUALIZATIONS

# VISUALIZATIONS

- We use ggplot2 and interactive ggvis as a standard, if possible
- Some plots use Viper Charts as backend (cost curves, lift charts, ...)
- Current GSOC project with Zach Jones
  - ▶ Demo plots for models in teaching
  - ▶ ROC curves
  - ▶ Threshold vs. Performance
  - ▶ Partial dependency plot
  - ▶ Learning curves

# R EXAMPLE

Visualizations

## Section 10

CARET VS. MLR

# CARET VS. MLR I

## WHY WE LIKE MLR MORE

- `mlr` has (in our opinion) a better OO design, class structure and infrastructure for future development and for users
- This makes things combinable, extensible, predictable
- More flexible parallelization with `parallelMap`, even on HPCs via `BatchJobs`
- Tuning with advanced methods like `irace`
- Fusing learners with other operations like pre-processing and feature selection
- Nested resampling (required for unbiased results)
- Survival and cluster analysis

## Section 11

## OPENML


# OPENML-R-PACKAGE I

Caution: Work in progress

OPENML?

- Main idea: Make ML experiments reproducible and most parts computer-readable
- Share everything
- Enrich with meta-information
- Later: Mine the results, meta-learn on it

## 1 minute intro




**Data** from  
various sources  
**analysed and  
organised online**  
for easy access

Scientists can **broadcast data**, explaining the challenge that needs to be addressed. OpenML will (for known data formats) **automatically analyze the data**, compute data characteristics, **annotate and index it for easy search**

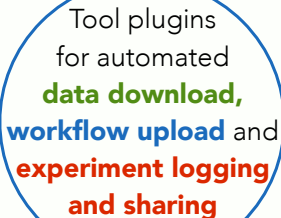


## 1 minute intro



Tasks are **realtime (collaborative) data mining challenges**, allowing anyone to build on previous results. OpenML creates **machine-readable descriptions** so that tools can **automatically download data**, use the correct procedures, and **upload all results online**.


## 1 minute intro



Tool plugins  
for automated  
**data download,**  
**workflow upload** and  
**experiment logging**  
**and sharing**

**Flows** are implementations of algorithms, workflows, or scripts **solving OpenML tasks**. OpenML keeps track of **flow details and versioning**, **organizes all their results** for easy comparison, even across tools.

## 1 minute intro



**Experiments**  
**auto-uploaded,**  
linked to **data, flows**  
and **authors**, and  
organised for easy  
reuse

**Runs** contain the results that **flows** obtained on specific tasks. Runs are **fully reproducible**, linked to the underlying data, tasks, flows and authors. OpenML **organizes all results online** for **discovery, comparison and reuse**

# OPENML-R-PACKAGE I

Let's visit website and project page

# OPENML-R-PACKAGE II

`https://github.com/openml/r`

## CURRENT API IN R

- Explore data and tasks
- Download data and tasks
- Register learners
- Upload runs
- Explore your own and other people's results

Already nicely connected to `mlr`!

# OPENML: EXPLORE AND SELECT DATA I

```
> library(OpenML)
> authenticateUser() # uses my OML config file

## Authenticating user at server: bernd_bischl@gmx.net
## Retrieved session hash. Valid until: 2015-06-30 11:37:13

> listOMLDataSets()[1:3, 1:5]

## Downloading 'http://www.openml.org/api/?f=openml.data' to '<mem>'

##   did status      name NumberOfClasses NumberOfFeatures
## 1  1  active    anneal                6             39
## 2  2  active anneal.ORIG                6             39
## 3  3  active   kr-vs-kp                 2             37
```

# OPENML: EXPLORE AND SELECT DATA II

```
> listOMLTasks()[1:3, 1:5]
```

```
## Downloading
```

```
'http://www.openml.org/api/?f=openml.tasks&task_type_id=1' to '<mem>'
```

##	task_id	task_type	did	status	name
## 1	1	Supervised Classification	1	active	anneal
## 2	2	Supervised Classification	2	active	anneal.ORIG
## 3	3	Supervised Classification	3	active	kr-vs-kp

# OPENML: DOWNLOAD A DATA SET

```
> # uses built in caching from disk
> getOMLDataSet(6)

## Getting data set '6' from OpenML repository.
## Downloading
'http://www.openml.org/api/?f=openml.data.description&data.id=6' to
'/tmp/RtmpHTqqZA/cache/datasets/6/description.xml'
## Downloading
'http://openml.liacs.nl/data/download/6/dataset_6_letter.arff' to
'/tmp/RtmpHTqqZA/cache/datasets/6/dataset.arff'

##
## Data Set "letter" :: (Version = 1, OpenML ID = 6)
## Default Target Attribute: class
```



# OPENML: DOWNLOAD A TASK I

```
> # uses built in caching from disk
> oml.task = getOMLTask(1)

## Downloading task '1' from OpenML repository.
## Downloading 'http://www.openml.org/api/?f=openml.task.get&task_id=1'
to '/tmp/Rtmp0QG2mN/cache/tasks/1/task.xml'
## Authenticating user at server: bernd_bischl@gmx.net
## Retrieved session hash. Valid until: 2015-06-30 12:19:28
## Getting data set '1' from OpenML repository.
## Downloading
'http://www.openml.org/api/?f=openml.data.description&data.id=1' to
'/tmp/Rtmp0QG2mN/cache/datasets/1/description.xml'
## Downloading
'http://openml.liacs.nl/data/download/1/dataset_1_anneal.arff' to
'/tmp/Rtmp0QG2mN/cache/datasets/1/dataset.arff'
## Downloading
'http://www.openml.org/api_splits/get/1/Task_1_splits.arff' to
'/tmp/Rtmp0QG2mN/cache/tasks/1/datasplits.arff'
```

# OPENML: DOWNLOAD A TASK II

```
> print(oml.task)

##
## OpenML Task 1 :: (Data ID = 1)
##   Task Type           : Supervised Classification
##   Data Set            : anneal :: (Version = 2, OpenML ID = 1)
##   Target Feature(s)   : class
##   Estimation Procedure : Stratified crossvalidation (1 x 10 folds)
```

# OPENML: RUN A TASK


```
> lrn = makeLearner("classif.rpart")
> res = runTaskMlr(oml.task, lrn)

## Warning in fixupData.Task(task, target, fixup.data): Empty factor levels were dropped for
columns:
family,product.type,steel,condition,formability,surface.finish,enamellability,m,marvi,corr,blue.2fbright.2
## Removing 7 columns: product.type,m,marvi,corr,jurofm,s,p
## [Resample] cross-validation iter: 1
## [Resample] cross-validation iter: 2
## [Resample] cross-validation iter: 3
## [Resample] cross-validation iter: 4
## [Resample] cross-validation iter: 5
## [Resample] cross-validation iter: 6
## [Resample] cross-validation iter: 7
## [Resample] cross-validation iter: 8
## [Resample] cross-validation iter: 9
## [Resample] cross-validation iter: 10
## [Resample] Result: acc.test.mean=0.977
## Downloading
'http://www.openml.org/api/?f=openml.implementation.exists&name=classif.rpart&external_version=R_7f351643'
to '<mem>'
```

# OPENML: UPLOAD LEARNER AND PREDICTIONS

```
> hash = authenticateUser("your@email.com", "your_password")
> impl = createOpenMLImplementationForMlrLearner(lrn)
> uploadOpenMLImplementation(impl, session.hash = hash)
> uploadOpenMLRun(oml.task, lrn, impl, pred, hash)
```

# Towards OpenML in education



frontier Joaquin Vanschoren  
Olav Bunte Stephan Oostveen  
Jorn Engelbart Mathijs van Lier  
Stefan Majoer



**Rogier Beckers**

@RogierBeckers



Follow

Het bewijs dat ik studeer op zondag!  
“@joavanschoren: #Machinelearning students  
on a #collaborative data mining ”


View translation

Lauradorp, Landgraaf



Contributions over time

every point is a task, click for details



frontier Joaquin Vanschoren Perry van Wesel Jose Melo Jos Mangnus Daan Peters Tom Becht Kevin Jacobs Koen Engelen  
Olav Bunte Stephan Oostveen Ray van den Hurk Schwester Kogowski Ky-Anh Tran Edgar Salas Thomas Tiel Groenesteghe  
Jorn Engelbart Mathijs van Lier Henry He Riche Brundensieck Hugo Lape Stanley Clark Christoforos Bousmoulas Rogier Beckers  
Stefan Majoer

RETWEETS

2

FAVORITES

2



9:48 PM - 7 Dec 2014



Kevin Jacobs Koen Engelen  
Tiel Groenesteghe  
Poukouvelas Rogier Beckers

Section 12

THE END

# THERE IS MORE ...

- Regular cost-sensitive learning (class-specific costs)
- Cost-sensitive learning (example-dependent costs)
- Model-based optimization
- Multi-criteria optimization
- OpenML
- ...

# OUTLOOK

## WE ARE WORKING ON

- Even better tuning system
- More interactive plots
- Large-Scale SVM ensembles
- Time-Series tasks
- Better benchmark analysis
- Multi-Label classification
- ...



Thanks!